

# The Count Distinct Problem

Steven Rosendahl

# The Question

- ▶ What is the problem?

# The Question

- ▶ What is the problem?
- ▶ Imagine we have a set called  $\mathbb{V}$  that contains a billion elements of the same type.

# The Question

- ▶ What is the problem?
- ▶ Imagine we have a set called  $\mathbb{V}$  that contains a billion elements of the same type.
- ▶ How many *unique* elements are in  $\mathbb{V}$ ?

# Questions

1. Pokémon Problem: How many unique Pokémon will a player encounter in a given playthrough of all the games?

# Questions

1. Pokémon Problem: How many unique Pokémon will a player encounter in a given playthrough of all the games?
2. Twitter Problem: How many unique hashtags are made a day on Twitter?

# Questions

1. Pokémon Problem: How many unique Pokémon will a player encounter in a given playthrough of all the games?
2. Twitter Problem: How many unique hashtags are made a day on Twitter?

We will use  $\mathbb{S}$  to represent the set of all the data, and  $\mathbb{V}$  to represent the set of unique elements.

# Hashing

- ▶ What is *hashing*?



# Hashing

- ▶ What is *hashing*?
- ▶ Applying a function  $h(x)$  to every element in  $\mathbb{S}$ , and storing the result in  $\mathbb{V}$ .

# Hashing

- ▶ What is *hashing*?
- ▶ Applying a function  $h(x)$  to every element in  $\mathbb{S}$ , and storing the result in  $\mathbb{V}$ .
- ▶ Ideally,  $h(x)$  is

# Hashing

- ▶ What is *hashing*?
- ▶ Applying a function  $h(x)$  to every element in  $\mathbb{S}$ , and storing the result in  $\mathbb{V}$ .
- ▶ Ideally,  $h(x)$  is
  1. Onto (surjective)

# Hashing

- ▶ What is *hashing*?
- ▶ Applying a function  $h(x)$  to every element in  $\mathbb{S}$ , and storing the result in  $\mathbb{V}$ .
- ▶ Ideally,  $h(x)$  is
  1. Onto (surjective)
  2. One-to-one (injective)

# Hashing

- ▶ What is *hashing*?
- ▶ Applying a function  $h(x)$  to every element in  $\mathbb{S}$ , and storing the result in  $\mathbb{V}$ .
- ▶ Ideally,  $h(x)$  is
  1. Onto (surjective)
  2. One-to-one (injective)
- ▶ We can ignore the duplicate values in  $\mathbb{V}$ .

# Solving the Pokémon Problem

- ▶ How can we solve the Pokémon Problem using a hash?

# Solving the Pokémon Problem

- ▶ How can we solve the Pokémon Problem using a hash?
- ▶ Create a hash function that turns a given Pokémon into a numerical value

# Solving the Pokémon Problem

- ▶ How can we solve the Pokémon Problem using a hash?
- ▶ Create a hash function that turns a given Pokémon into a numerical value
- ▶ Store the result in  $\mathbb{V}$  if it is not already there.



# Solving the Pokémon Problem

- ▶ How can we solve the Pokémon Problem using a hash?
- ▶ Create a hash function that turns a given Pokémon into a numerical value
- ▶ Store the result in  $\mathbb{V}$  if it is not already there.
- ▶ The hash function:
  1. Sum up the ASCII value of each character in a Pokémon's name. Call this  $n$ .

# Solving the Pokémon Problem

- ▶ How can we solve the Pokémon Problem using a hash?
- ▶ Create a hash function that turns a given Pokémon into a numerical value
- ▶ Store the result in  $\mathbb{V}$  if it is not already there.
- ▶ The hash function:
  1. Sum up the ASCII value of each character in a Pokémon's name. Call this  $n$ .
  2. Add  $n$  to the Pokémon's corresponding National Pokédex number. Call this  $m$ .

# Solving the Pokémon Problem

- ▶ How can we solve the Pokémon Problem using a hash?
- ▶ Create a hash function that turns a given Pokémon into a numerical value
- ▶ Store the result in  $\mathbb{V}$  if it is not already there.
- ▶ The hash function:
  1. Sum up the ASCII value of each character in a Pokémon's name. Call this  $n$ .
  2. Add  $n$  to the Pokémon's corresponding National Pokédex number. Call this  $m$ .
  3. Find  $m \bmod 721$ .

# Solving the Pokémon Problem

- ▶ How can we solve the Pokémon Problem using a hash?
- ▶ Create a hash function that turns a given Pokémon into a numerical value
- ▶ Store the result in  $\mathbb{V}$  if it is not already there.
- ▶ The hash function:
  1. Sum up the ASCII value of each character in a Pokémon's name. Call this  $n$ .
  2. Add  $n$  to the Pokémon's corresponding National Pokédex number. Call this  $m$ .
  3. Find  $m \bmod 721$ .
  4. Results:

# Solving the Pokémon Problem

- ▶ How can we solve the Pokémon Problem using a hash?
- ▶ Create a hash function that turns a given Pokémon into a numerical value
- ▶ Store the result in  $\mathbb{V}$  if it is not already there.
- ▶ The hash function:
  1. Sum up the ASCII value of each character in a Pokémon's name. Call this  $n$ .
  2. Add  $n$  to the Pokémon's corresponding National Pokédex number. Call this  $m$ .
  3. Find  $m \bmod 721$ .
  4. Results:
    - ▶ On average: 461 unique encounters

# Implementation

```
1 array.each do |pokemon|
2     hash_table[pokemon.hash] = pokemon.hash
3 end
4 total = 0
5 hash_table.each do |val|
6     if val != -1
7         total = total + 1
8     end
9 end
10 puts "The total number of unique encounters is #{total}"
```

# Problems With The Hash Table

- ▶ Memory Intensive
  - ▶ Pokémon problem only dealt with a set  $\mathbb{S}$  of size 6000

# Problems With The Hash Table

- ▶ Memory Intensive
  - ▶ Pokémon problem only dealt with a set  $\mathbb{S}$  of size 6000
  - ▶ Twitter Problem deals with  $\mathbb{S}$  of size 200,000,000.



# Problems With The Hash Table

- ▶ Memory Intensive
  - ▶ Pokémon problem only dealt with a set  $\mathbb{S}$  of size 6000
  - ▶ Twitter Problem deals with  $\mathbb{S}$  of size 200,000,000.
  - ▶ Collisions and collision policies also add to the amount of memory required.

# The Algorithm

1. Create a bitmap in memory. We will call this  $\mathbb{V}$ .

# The Algorithm

1. Create a bitmap in memory. We will call this  $\mathbb{V}$ .
2. For each value  $s$  in  $\mathbb{S}$ , hash  $s$  to a binary number.

# The Algorithm

1. Create a bitmap in memory. We will call this  $\mathbb{V}$ .
2. For each value  $s$  in  $\mathbb{S}$ , hash  $s$  to a binary number.
  - ▶ We will use a *Murmur Hash* to do this.

# The Algorithm

1. Create a bitmap in memory. We will call this  $\mathbb{V}$ .
2. For each value  $s$  in  $\mathbb{S}$ , hash  $s$  to a binary number.
  - ▶ We will use a *Murmur Hash* to do this.
3. Analyze the first sequence of 0's in the binary value, and store the number of leading 0's into the bitmap.
  - ▶ If the sequence of 0's has been seen before, there is a high probability that the term is a duplicate

# The Algorithm

1. Create a bitmap in memory. We will call this  $\mathbb{V}$ .
2. For each value  $s$  in  $\mathbb{S}$ , hash  $s$  to a binary number.
  - ▶ We will use a *Murmur Hash* to do this.
3. Analyze the first sequence of 0's in the binary value, and store the number of leading 0's into the bitmap.
  - ▶ If the sequence of 0's has been seen before, there is a high probability that the term is a duplicate
4. Take the harmonic average of all the totals in the bitmap.

# The Math

- ▶ How much memory is required?

# The Math

- ▶ How much memory is required?

$$\text{Memory Required} = \log_2 (\log_2 (M))$$



# The Math

- ▶ How much memory is required?

$$\text{Memory Required} = \log_2 (\log_2 (M))$$

- ▶  $M$  is the size of the original set of data ( $\mathbb{S}$ ).

# The Math

- ▶ How much memory is required?

$$\text{Memory Required} = \log_2 (\log_2 (M))$$

- ▶  $M$  is the size of the original set of data ( $\$$ ).
- ▶ For the Twitter Problem:

# The Math

- ▶ How much memory is required?

$$\text{Memory Required} = \log_2 (\log_2 (M))$$

- ▶  $M$  is the size of the original set of data ( $\$$ ).
- ▶ For the Twitter Problem:

$$\begin{aligned}\text{Memory Required} &\approx \log_2 (\log_2 (200,000,000 \times 10)) \\ &\approx 4.94kb\end{aligned}$$

# The Math

- ▶ How much memory is required?

$$\text{Memory Required} = \log_2 (\log_2 (M))$$

- ▶  $M$  is the size of the original set of data ( $\$$ ).
- ▶ For the Twitter Problem:

$$\begin{aligned}\text{Memory Required} &\approx \log_2 (\log_2 (200,000,000 \times 10)) \\ &\approx 4.94kb\end{aligned}$$

- ▶ What is the predicted error?

# The Math

- ▶ How much memory is required?

$$\text{Memory Required} = \log_2 (\log_2 (M))$$

- ▶  $M$  is the size of the original set of data ( $\mathbb{S}$ ).
- ▶ For the Twitter Problem:

$$\begin{aligned}\text{Memory Required} &\approx \log_2 (\log_2 (200,000,000 \times 10)) \\ &\approx 4.94kb\end{aligned}$$

- ▶ What is the predicted error?

$$\text{Error} = \frac{\sqrt{3 \log(2) - 1}}{m}$$

# The Math

- ▶ How much memory is required?

$$\text{Memory Required} = \log_2 (\log_2 (M))$$

- ▶  $M$  is the size of the original set of data ( $\mathbb{S}$ ).
- ▶ For the Twitter Problem:

$$\begin{aligned}\text{Memory Required} &\approx \log_2 (\log_2 (200,000,000 \times 10)) \\ &\approx 4.94kb\end{aligned}$$

- ▶ What is the predicted error?

$$\text{Error} = \frac{\sqrt{3 \log(2) - 1}}{m}$$

- ▶  $m$  is the number of spaces in the bitmap ( $\mathbb{V}$ ).

# Solving the Twitter Problem

- ▶ How can we process count 200,000,000 hashtags on an average computer?

# Solving the Twitter Problem

- ▶ How can we process count 200,000,000 hashtags on an average computer?
- ▶ We can lower the sample size and apply a best fit line to the data.



# Solving the Twitter Problem

- ▶ How can we process count 200,000,000 hashtags on an average computer?
- ▶ We can lower the sample size and apply a best fit line to the data.
  1. For 24 hours, gather 2000 tweets containing “#” every 2 minutes.

# Solving the Twitter Problem

- ▶ How can we process count 200,000,000 hashtags on an average computer?
- ▶ We can lower the sample size and apply a best fit line to the data.
  1. For 24 hours, gather 2000 tweets containing “#” every 2 minutes.
  2. Using the HyperLogLog, determine the unique number of total hashtags every time a new sample is gathered.

# Implementation

```
1  mhll = Hyperll::HyperLogLog.new(10)
2  File.open("twitter_data.txt", "r") do |file|
3      file.each_line do |line|
4          mhll.offer line
5      end
6  end
7  str = "Unique Elements: #{mhll.cardinality}"
8  puts str
```

# Results




Figure:  $0.284356x + 7361.39$

# Results




Figure:  $0.284356x + 7361.39$

- Plugging in 200,000,000 gives us

# Results

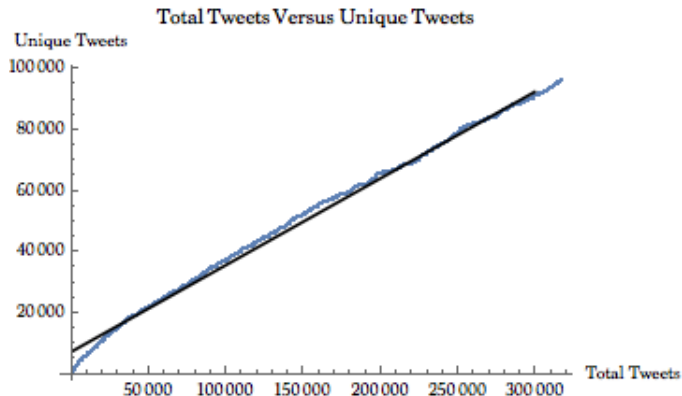


Figure:  $0.284356x + 7361.39$

- ▶ Plugging in 200,000,000 gives us
- ▶  $5.68785 \times 10^7$  unique hashtags.